

Parallel Computing with Bayesian MCMC and Data Cloning in R with the `dclone` package

Péter Sólymos
University of Alberta

Abstract

The `dclone` R package provides parallel computing features for Bayesian MCMC computations and the use of the data cloning algorithm. Parallelization can be achieved by running independent parallel MCMC chains or by partitioning the problem into subsets. Size balancing is introduced as a simple workload optimization algorithm applicable when processing time for the pieces differ significantly. The speed-up with parallel computing infrastructure can facilitate the analysis of larger data and complex hierarchical models building upon commonly available Bayesian software.

Keywords: Bayesian statistics, data cloning, maximum likelihood inference, generalized linear models, R, parallel computing.

1. Introduction

Large data and complex hierarchical models pose a challenge to statistical computations. For this end researchers often use parallel computing, especially for “embarrassingly parallel problems” such as Markov chain Monte Carlo (MCMC) methods. R (R Development Core Team 2010) is an open source software and environment for statistical computing and has interfaces for most commonly used Bayesian software (JAGS, Plummer (2010); WinBUGS, Spiegelhalter *et al.* (2003); and OpenBUGS Spiegelhalter *et al.* (2007)). R also has capabilities to exploit capacity of multi-core machines or idle computers in a network through various contributed packages (Schmidberger *et al.* 2009). R is thus a suitable platform for parallelizing Bayesian MCMC computations, which is not a built-in feature in commonly used Bayesian software.

Data cloning is a global optimization algorithm that exploits Markov chain Monte Carlo (MCMC) methods used in the Bayesian statistical framework while providing valid frequentist inferences such as the maximum likelihood estimates and their standard errors (Lele *et al.* 2007, 2010). This approach enables to fit complex hierarchical models (Lele *et al.* 2007; Ponciano *et al.* 2009) and helps in studying parameter identifiability issues (Lele *et al.* 2010; Sólymos 2010), but comes with a price in processing time that increases with the number of clones by repeating the same data many times, thus increasing graph dimensions in the Bayesian software used. Again, parallelization naturally follows as a reply to the often demanding computational requirements of data cloning.

The R package `dclone` (Sólymos 2010) provide convenience functions for Bayesian computations and data cloning. The package also has functions for parallel computations, building on

the parallel computing infrastructure provided by the **snw** package (Tierney *et al.* 2010). In this paper I present best practices for parallel Bayesian computations in R, describe how the wrapper functions in the **dclone** package work, demonstrate how computing time can improve by these functions and how parallelization can be effectively optimized.

2. Bayesian example

We will use the **seeds** example from the WinBUGS Manual Vol. I. (Spiegelhalter *et al.* 2003) based on Table 3 of Crowder (1978). The experiment concerns the proportion of seeds that germinated on each of 21 plates arranged according to a 2×2 factorial layout by seed (x_{1i}) and type of root extract (x_{2i}). The data are the number of germinated (r_i) and the total number of seeds (n_i) on the i th plate, $i = 1, \dots, N$:

```
R> source("http://dcr.r-forge.r-project.org/examples/seeds.R")
R> str(dat <- seeds$data)
```

```
List of 5
 $ N : num 21
 $ r : num [1:21] 10 23 23 26 17 5 53 55 32 46 ...
 $ n : num [1:21] 39 62 81 51 39 6 74 72 51 79 ...
 $ x1: num [1:21] 0 0 0 0 0 0 0 0 0 0 ...
 $ x2: num [1:21] 0 0 0 0 0 1 1 1 1 1 ...
```

The Bayesian model is random effects logistic, allowing for overdispersion:

$$\begin{aligned} r_i &\sim \text{Binomial}(p_i, n_i) \\ \text{logit}(p_i) &= \alpha_0 + \alpha_1 x_{1i} + \alpha_2 x_{2i} + \alpha_{12} x_{1i} x_{2i} + b_i \\ b_i &\sim \text{Normal}(0, \sigma^2), \end{aligned}$$

where p_i is the probability of germination on the i th plate, and $\alpha_0, \alpha_1, \alpha_2, \alpha_{12}$ and precision parameter $\tau = 1/\sigma^2$ are given independent noninformative priors. The corresponding BUGS model is:

```
R> (model <- seeds$model)

function() {
  alpha0 ~ dnorm(0.0, 1.0E-6);      # intercept
  alpha1 ~ dnorm(0.0, 1.0E-6);      # seed coeff
  alpha2 ~ dnorm(0.0, 1.0E-6);      # extract coeff
  alpha12 ~ dnorm(0.0, 1.0E-6);     # intercept
  tau ~ dgamma(1.0E-3, 1.0E-3);    # 1/sigma^2
  sigma <- 1.0/sqrt(tau);
  for (i in 1:N) {
    b[i] ~ dnorm(0.0, tau);
    logit(p[i]) <- alpha0 + alpha1*x1[i] + alpha2*x2[i] +
      alpha12*x1[i]*x2[i] + b[i];
    r[i] ~ dbin(p[i], n[i]);
  }
}
```

Initial values are taken from the WinBUGS manual

```
R> str(inits <- seeds$inits)
```

```
List of 5
```

```
$ tau      : num 1
$ alpha0   : num 0
$ alpha1   : num 0
$ alpha2   : num 0
$ alpha12  : num 0
```

We are using the following settings (`n.adapt` + `n.update` = burn-in iterations, `n.iter` is the number of samples taken from the posterior distribution, `thin` is the thinning value, `n.chains` is the number of chains used) for fitting the JAGS model and monitoring the parameters in `params`:

```
R> n.adapt <- 1000
R> n.update <- 1000
R> n.iter <- 3000
R> thin <- 10
R> n.chains <- 3
R> params <- c("alpha0", "alpha1", "alpha2", "alpha12", "sigma")
```

The full Bayesian results can be obtained via JAGS as

```
R> m <- jags.fit(data = dat, params = params, model = model,
+   inits = inits, n.adapt = n.adapt, n.update = n.update,
+   n.iter = n.iter, thin = thin, n.chains = n.chains)
```

3. Parallel MCMC chains

Markov chain Monte Carlo (MCMC) methods represent an “embarrassingly parallel problem” (Rossini *et al.* 2003; Schmidberger *et al.* 2009) and efforts are being made to enable effective parallel computing techniques for Bayesian statistics (Wilkinson 2005), but popular software for fitting graphical models using some dialect of the BUGS language (e.g. WinBUGS, OpenBUGS, JAGS) do not yet come with a built in facility for enabling parallel computations (see critique for Lunn *et al.* (2009)).

WinBUGS has been criticized that the independence of parallel chains cannot be guaranteed via setting random number generator (RNG) seed (Wilkinson 2005). Consequently, parallelizing approaches built on WinBUGS (e.g. GridBUGS Girgis *et al.* (2007); or bugsparell Metrum Institute (2010)) are not best suited for all types of parallel Bayesian computations (e.g. those are best suited for batch processing, but not for generating truly independent chains). One might argue that the return times of RNG algorithms are large enough to be safe to use a *different seeds* approach, but this approach might be risky for long chains as parallel streams might have some overlap, and parts of the sequences might not be independent (see Wilkinson (2005) for discussion). Also, using proper parallel RNG of the `rlecuyer` package (Sevcikova and Rossini 2009) used in the R session to ensure independence of initial values or random seeds will not guarantee chain independence outside of R (e.g. this is the case in the `bugsparell` project according to its source code).

Random numbers for initialization in JAGS are defined by the type and seed of the RNG, and currently there are four different types of RNGs available (`base::Wichmann-Hill`, `base::-`

Marsaglia-Multicarry, base::Super-Duper, base::Mersenne-Twister). The model initialization by default uses different RNG types for up to four chains, which naturally leads to safe parallelization because these can represent independent streams on each processor.

Now let us fit the model by using MCMC chains computed on parallel workers. We use three workers (one worker for each MCMC chain, the worker type we use here is "socket", but type can be anything else accepted by the `snow` package, such as "PVM", "MPI", and "NWS"):

```
R> library(snow)
R> cl <- makeCluster(3, type = "SOCK")
```

Because MCMC computations (except for generating user supplied initial values) are done outside of R, we just have to ensure that model initialization is done properly and initial values contain different `.RNG.state` and `.RNG.name` values. Up to four chains (JAGS has 4 RNG types) this can be most easily achieved by

```
R> inits2 <- jags.fit(dat, params, model, inits, n.chains,
+                   n.adapt = 0, n.update = 0, n.iter = 0)$state(internal = TRUE)
```

Note that this initialization might represent some overhead especially for models with many latent nodes, but this is the most secure way of ensuring the independence of MCMC chains.

We have to load the `dclone` package on the workers and set a common working directory where the model file can be found (a single model file is used to avoid multiple read/write requests):

```
R> clusterEvalQ(cl, library(dclone))
R> clusterEvalQ(cl, setwd(getwd()))
R> filename <- write.jags.model(model)
```

Using a common file is preferred when the memory is shared (e.g. for `SOCK` cluster type). If computations are made on computing grid or cluster without shared memory, each machine needs a model file written onto the hard drive. For this, one can use the `write.jags.model` function with argument `overwrite = TRUE` to ensure all worker has the same model file name. In this case, the name of the model file can be provided as a character vector in subsequent steps.

The next step is to distribute the data to the workers

```
R> cldata <- list(data=dat, params=params, model=filename, inits=inits2)
R> clusterExport(cl, "cldata")
```

and create a function that will be evaluated on each worker

```
R> jagsparallel <- function(i, ...) {
+   jags.fit(data = cldata$data, params = cldata$params, model = cldata$model,
+           inits = cldata$inits[[i]], n.chains = 1, updated.model = FALSE, ...)
+ }
```

Now we can use the `parLapply` function of the `snow` package and then clean up the model file from the hard drive

```
R> res <- parLapply(cl, 1:n.chains, jagsparallel,
+                 n.adapt = n.adapt, n.update = n.update, n.iter = n.iter, thin = thin)
R> res <- as.mcmc.list(lapply(res, as.mcmc))

R> clean.jags.model(filename)
```

The `jags.parfit` function in the **dclone** package is a wrapper for this process, here we can use the `model` function instead of a model file

```
R> m2 <- jags.parfit(cl, data = dat, params = params, model = model,
+               inits = inits, n.adapt = n.adapt, n.update = n.update,
+               n.iter = n.iter, thin = thin, n.chains = n.chains)
R> stopCluster(cl)
```

The `jags.parfit` is the parallel counterpart of the `jags.fit` function. It internally uses the `snowWrapper` function because `clusterExport` exports objects from the global environment and not from its child environments (i.e. from inside the `jags.parfit` function). The `snowWrapper` function takes care of name clashes and also can be used to set the parallel random number generation type via the `clusterSetupRNG` function of **snow**. The **dclone** package can use the L'Ecuyer's RNG (L'Ecuyer *et al.* 2002) ("RNGstream", requires the **rlecuyer** package, Sevcikova and Rossini (2009)) or "SPRNG" (requires the **rsprng** package, Li (2010)). The type of parallel RNG can be set up by the `dcoptions` function and by default it is "none" (not setting up RNG on the workers). Setting the parallel RNG can be important when using random numbers as initial values supplied as function, so the function is evaluated on the workers. To set the RNG type, we use

```
R> dcoptions(RNG = "RNGstream")
R> dcoptions()$RNG
```

```
[1] "RNGstream"
```

Table 1 shows different ways of providing initial values. Besides setting up RNGs, we might have to export objects to the workers using the `clusterExport` function of the **snow** package if `inits` contains a reference to objects defined in the global environment.

4. Parallel data cloning

Data cloning uses MCMC methods to calculate the posterior distribution of the model parameters conditional on the data. If the number of clones (k) is large enough, the mean of the posterior distribution is the maximum likelihood estimate and k times the posterior variance is the corresponding asymptotic variance of the maximum likelihood estimate (Lele *et al.* 2007, 2010). But there is no universal number of clones that is needed to reach convergence, so the number of clones has to be determined empirically by using e.g. the `dctable` and `dcdiag` functions provided by the **dclone** package (Sólymos 2010). For this, we have to fit the same model with the different number of clones of the data, and check if the variances of the parameters are going to zero (Lele *et al.* 2010). Fitting the model with number of clones, say (1, 2, 5, 10, 25), means that we have to fit the model five times. Without parallelization, we can use subsequent calls of the `jags.fit` function with different data sets corresponding to different number of clones, or use the `dc.fit` wrapper function of the **dclone** R package.

Parallelization can happen in two ways: (1) we parallelize the computation of the MCMC chains (usually we use more than one chain to check proper mixing behaviour of the chains); or (2) we split the problem into subsets and run the subsets on different worker of the parallel computing environment. The first type of parallelization can be addressed by multiple calls to the `jags.parfit` function. The second type of parallelization is implemented in the `dc.parfit` function that is the parallelized version of the function `dc.fit`.

<code>inits</code>	Effect	Notes for parallel execution
NULL	Initial values automatically generated with RNG type and seed	None
Self contained named list	<code>inits</code> used, RNG type and seed generated if missing	None
Named list with references to global environment	<i>ditto</i>	Objects must be exported by <code>clusterExport()</code>
Self contained function	Function used to generate <code>inits</code> , RNG type and seed generated if missing	None
Function using random numbers	<i>ditto</i>	Parallel RNG should be set by <code>dcoptions()</code> or <code>clusterSetupRNG()</code>
Fuction with references to global environment	<i>ditto</i>	Objects must be exported by <code>clusterExport()</code>

Table 1: Different ways of providing initial values and notes for parallel executions. Self contained means that the object has no reference to other objects or environments, neither random numbers are generated within a self contained function.

4.1. Data cloning with parallel MCMC chains

The following simple function allows us to switch parallelization on/off (by setting the `parallel` argument as `TRUE/FALSE`, respectively), and also measure the processing time in minutes:

```
R> timerfitfun <- function(parallel = FALSE, ...) {
+   t0 <- proc.time()
+   mod <- if (parallel)
+     jags.parfit(...) else jags.fit(...)
+   attr(mod, "timer") <- (proc.time() - t0)[3] / 60
+   mod
+ }
```

The vector `n.clones` contains the number of clones to be used:

```
R> n.clones <- c(1, 2, 5, 10, 25)
```

First we fit the five models sequentially (one model for each element of `n.clones`) without parallelization of the three MCMC chains:

```
R> res1 <- lapply(n.clones, function(z)
+   timerfitfun(parallel = FALSE,
+     data = dclone(dat, n.clones= z, multiply = "N"),
+     params = params,
+     model = model,
+     inits = inits,
+     n.adapt = n.adapt,
+     n.update = n.update,
```

```
+      n.iter = n.iter,
+      thin = thin,
+      n.chains = n.chains))
```

Now we fit the models by using MCMC chains computed on parallel workers. We use three workers, one worker for each MCMC chain:

```
R> cl <- makeCluster(3, type = "SOCK")
R> res2 <- lapply(n.clones, function(z)
+   timerfitfun(parallel = TRUE,
+   cl = cl,
+   data = dclone(dat, n.clones = z, multiply = "N"),
+   params = params,
+   model = model,
+   inits = inits,
+   n.adapt = n.adapt,
+   n.update = n.update,
+   n.iter = n.iter,
+   thin = thin,
+   n.chains = n.chains))
R> stopCluster(cl)
```

We can extract timer information from the result objects¹:

```
R> pt1 <- sapply(res1, function(z) attr(z, "timer"))
R> pt2 <- sapply(res2, function(z) attr(z, "timer"))
```

Processing time (column `time`) increased linearly with `n.clones` (with and without parallelization), columns `rel.change` indicate change relative to `n.clones = 1`, `speedup` is the speed-up of the parallel computation relative to its sequential counterpart:

```
R> tab1 <- data.frame(n.clones=n.clones,
+   time=pt1,
+   rel.change = pt1 / pt1[1],
+   time.par=pt2,
+   rel.change.par = pt2 / pt2[1],
+   speedup = pt2 / pt1)
R> round(tab1, 2)
```

	n.clones	time	rel.change	time.par	rel.change.par	speedup
1	1	0.16	1.00	0.05	1.00	0.34
2	2	0.28	1.74	0.09	1.61	0.31
3	5	0.67	4.23	0.25	4.58	0.37
4	10	1.35	8.52	0.44	8.23	0.33
5	25	3.32	20.94	1.16	21.48	0.35

It took a total of 5.77 minutes to fit the five models without, and only 1.99 minutes (34.4 %) with parallel MCMC chains. Parallel MCMC computations have effectively reduced the processing time almost at a rate of $1/(\text{number of chains})$.

During iterative model fitting for data cloning using parallel MCMC chains, it is possible to use the joint posterior distribution from the previous iteration as a prior distribution for

¹All computations presented here were run on a computer with Intel® Core™ i7-920 CPU, 3 GB RAM and Windows XP operating system.

the next iteration (possibly defined as a Multivariate Normal prior, see [Sólymos \(2010\)](#) for implementation).

4.2. Partitioning into parallel subsets

The other way of parallelizing the iterative fitting procedure with data cloning is to split the problem into subsets. The `snow` package has the `clusterSplit` function to split a problem into equal sized subsets, and it is then used in e.g. the `parLapply` function. In this approach, each piece is assumed to be of equal size. The function would split the `n.clones` vector on two workers as

```
R> clusterSplit(1:2, n.clones)

[[1]]
[1] 1 2 5

[[2]]
[1] 10 25
```

But we have already learnt that running the first subset would take 1.1 while the second would take 4.67 minutes, thus the speed-up would be only 0.81. We have also learnt that computing time scaled linearly with the number of clones, thus we can use this information to optimize the work load.

The `clusterSplitSB` function of the `dclone` package implements *size balancing* that splits the problem into subsets with respect to size (i.e. approximate processing time). The function `parLapplySB` uses these subsets similarly to `parLapply`. Our `n.clones` vector is partitioned with two workers as

```
R> clusterSplitSB(1:2, n.clones, size = n.clones)

[[1]]
[1] 25

[[2]]
[1] 10 5 2 1
```

The expected speed-up based on results in `pt1` is 0.57.

In size balancing, the problems are re-ordered from largest to smallest, and then subsets are determined by minimizing the total approximate processing time. This splitting is deterministic, thus computations are reproducible. However, size balancing can be combined with load balancing ([Rossini et al. 2003](#)). This option is implemented in the `parLapplySLB` function. Load balancing is useful for utilizing inhomogeneous clusters in which nodes differ significantly in performance. With load balancing, the jobs for the initial list elements are placed on the processors, and the remaining jobs are assigned to the first available processor until all jobs are complete. If size (processing time) is correct, size plus load balancing should be identical to size balancing, but the actual sequence of execution might depend on unequal performance of the workers. The splitting in this case is non-deterministic (might not be reproducible). If the performance of the workers is known to differ significantly it is possible to switch to load balancing by setting `dcoptions(LB = TRUE)`, but it is turned off (`FALSE`) by default.

The `clusterSize` function can be used to determine the optimal number of workers needed for an (approximate) size vector:




Figure 1: Effect of balancing type on approximate total processing time by showing the subsets on each worker. Note, the actual sequence of execution might be different for load balancing, because it is non-deterministic. For data cloning, relative performance of balancing options can be well approximated by using the vector for the number of clones to be used.

```
R> clusterSize(n.clones)
```

	workers	none	load	size	both
1	1	43	43	43	43
2	2	35	31	25	25
3	3	35	27	25	25
4	4	25	26	25	25
5	5	25	25	25	25

The function returns approximate processing times (based on the size vector) as a function of the number of workers and balancing type. As we can see, only two workers are needed to accomplish the task by size balancing. "none" refers to splitting the problem into subsets without load or size balancing, "both" refers to load and size balancing, but it is identical to "size" balancing as improvement due to load balancing cannot be determined *a priori*.

The `plotClusterSize` function can help to visualize the effect of using different balancing options:

```
R> opar <- par(mfrow=c(1,3), cex.lab=1.5, cex.main=1.5, cex.sub=1.5)
R> col <- heat.colors(length(n.clones))
R> xlim <- c(0, max(clusterSize(n.clones)[2,-1]))
R> plotClusterSize(2, n.clones, "none", col = col, xlim = xlim)
R> plotClusterSize(2, n.clones, "load", col = col, xlim = xlim)
R> plotClusterSize(2, n.clones, "size", col = col, xlim = xlim)
R> par(opar)
```

We can see in Fig. 1, that, for two workers, size balancing results in the shortest processing time. While the model with `max(n.clones)` is running on one worker, all the other models can be fitted on the other worker. Total processing time is determined by the largest problem size.

First, we fit the five models iteratively without parallelization by using the `dc.fit` function of the `dclone` package:

```

R> t0 <- proc.time()
R> res3 <- dc.fit(dat, params, model, inits,
+               n.clones = n.clones, multiply = "N",
+               n.adapt = n.adapt, n.update = n.update,
+               n.iter = n.iter, thin = thin, n.chains = n.chains)
R> attr(res3, "timer") <- (proc.time() - t0)[3] / 60

```

Now let us fit the same models with parallelization using “socket” cluster type with two workers. The default balancing option in `dc.parfit` is size balancing. The function collects only descriptive statistics of the MCMC objects when number of clones is smaller than the maximum, similarly to `dc.fit`. The whole MCMC object is returned for the MCMC object with the largest number of clones. This approach reduces communication overhead.

```

R> cl <- makeCluster(2, type = "SOCK")
R> t0 <- proc.time()
R> res4 <- dc.parfit(cl, dat, params, model, inits,
+                  n.clones = n.clones, multiply = "N",
+                  n.adapt = n.adapt, n.update = n.update,
+                  n.iter = n.iter, thin = thin, n.chains = n.chains)
R> attr(res4, "timer") <- (proc.time() - t0)[3] / 60
R> stopCluster(cl)

```

Note, that by using the parallel partitioning approach to data cloning, it is not possible to use the posterior distribution as a prior in the next iteration with higher number of clones. It is only possible to use the same initial values for all runs, as opposed to the sequential `dc.fit` function (see [Sólymos \(2010\)](#) and `help("dc.fit")` for a description of the arguments `update`, `updatefun`, and `initsfun`). But opposed to the `jags.parfit` function which can work only with JAGS, the `dc.parfit` function can work with WinBUGS and OpenBUGS as well besides JAGS (see documentation for arguments `flavour` and `program` and examples on help pages `help("dc.parfit")` and `help("dc.fit")`).

It took 5.74 minutes without parallelization, as compared to 3.3 minutes (57.54 %) with parallelization. The parallel processing time in this case was close to the time needed to fit the model with `max(n.clones) = 25` clones without parallelization, so it was really the largest problem that determined the processing time.

5. Conclusions

The paper described the parallel computing features provided by the `dclone` R package for Bayesian MCMC computations and the use of the data cloning algorithm.

Large amount of parallelism can be achieved by many ways, e.g. parallelizing single chain ([Wilkinson 2005](#)) and via graphical processing units ([Suchard and Rambaut 2009](#); [Lee et al. 2010](#)). The parallel chains approach can be highly effective for Markov chains with good mixing and reasonably low burn-in times. The `jags.parfit` function fits models with parallel chains on separate workers by safely initializing the models, including RNG setup. Parallel MCMC chain computations are currently implemented to work with JAGS.

Size balancing is introduced as a simple workload optimization algorithm applicable when processing time for the pieces differ significantly. The package contains utility functions that enable size balancing and help optimizing the workload. The `dc.parfit` function uses size

balancing to iteratively fit models on multiple workers for data cloning. `dc.parfit` can work with JAGS, WinBUGS or OpenBUGS.

The functions presented here can facilitate the analysis of larger data and complex hierarchical models building upon commonly available Bayesian software.

6. Acknowledgments

I am grateful to Subhash Lele and Khurram Nadeem for valuable discussions. Comments from Christian Robert and the Associate Editor greatly improved the manuscript.

References

- Crowder M (1978). “Beta-Binomial ANOVA for proportions.” *Applied Statistics*, **27**, 34–37.
- Girgis IG, Schaible T, Nandy P, De Ridder F, Mathers J, Mohanty S (2007). *Parallel Bayesian methodology for population analysis*. URL <http://www.page-meeting.org/?abstract=1105>.
- L’Ecuyer P, Simard R, Chen EJ, Kelton WD (2002). “An Object-Oriented Random-Number Package With Many Long Streams and Substreams.” *Operations Research*, **50**(6), 1073–1075.
- Lee A, Yau C, Giles M, Doucet A, Holmes C (2010). “On the utility of graphics cards to perform massively parallel simulation of advanced Monte Carlo methods.” *Journal of Computational and Graphical Statistics*, **19**(4), 769–789.
- Lele SR, Dennis B, Lutscher F (2007). “Data cloning: easy maximum likelihood estimation for complex ecological models using Bayesian Markov chain Monte Carlo methods.” *Ecology Letters*, **10**, 551–563.
- Lele SR, Nadeem K, Schmuland B (2010). “Estimability and likelihood inference for generalized linear mixed models using data cloning.” *Journal of the American Statistical Association*, **105**, 1617–1625. In press.
- Li N (2010). *rsprng: R interface to SPRNG (Scalable Parallel Random Number Generators)*. R package version 1.0, URL <http://CRAN.R-project.org/package=rsprng>.
- Lunn D, Spiegelhalter D, Thomas A, Best N (2009). “The BUGS project: Evolution, critique and future directions.” *Statistics in Medicine*, **28**, 3049–3067. With discussion.
- Metrum Institute (2010). *bugsParallel: R scripts for parallel computation of multiple MCMC chains with WinBUGS*. URL <http://code.google.com/p/bugsparell/>.
- Plummer M (2010). *JAGS Version 2.2.0 manual (November 7, 2010)*. URL <http://mcmc-jags.sourceforge.net>.
- Ponciano JM, Taper ML, Dennis B, Lele SR (2009). “Hierarchical models in ecology: confidence intervals, hypothesis testing, and model selection using data cloning.” *Ecology*, **90**, 356–362.

- R Development Core Team (2010). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria. ISBN 3-900051-07-0, URL <http://www.R-project.org/>.
- Rossini A, Tierney L, Li N (2003). “Simple Parallel Statistical Computing in R.” *UW Biostatistics Working Paper Series*, **Working Paper 193.**, 1–27. URL <http://www.bepress.com/uwbiostat/paper193>.
- Schmidberger M, Morgan M, Eddelbuettel D, Yu H, Tierney L, Mansmann U (2009). “State of the Art in Parallel Computing with R.” *Journal of Statistical Software*, **31**(1), 1–27. ISSN 1548-7660. URL <http://www.jstatsoft.org/v31/i01>.
- Sevcikova H, Rossini T (2009). *rlecuyer: R interface to RNG with multiple streams*. R package version 0.3-1, URL <http://CRAN.R-project.org/package=rlecuyer>.
- Sólymos P (2010). “dclone: Data Cloning in R.” *The R Journal*, **2**(2), 29–37. R package version: 1.3-1, URL <http://journal.r-project.org/>.
- Spiegelhalter D, Thomas A, Best N, Lunn D (2007). *OpenBUGS User Manual, Version 3.0.2, September 2007*. URL <http://mathstat.helsinki.fi/openbugs/>.
- Spiegelhalter DJ, Thomas A, Best NG, Lunn D (2003). “WinBUGS Version 1.4 Users Manual.” MRC Biostatistics Unit, Cambridge. URL <http://www.mrc-bsu.cam.ac.uk/bugs/>.
- Suchard MA, Rambaut A (2009). “Many-Core Algorithms for Statistical Phylogenetics.” *Bioinformatics*, **25**(11), 1370–1376.
- Tierney L, Rossini AJ, Li N, Sevcikova H (2010). *snow: Simple Network of Workstations*. R package version 0.3-3, URL <http://cran.r-project.org/packages=snow>.
- Wilkinson DJ (2005). “Parallel Bayesian Computation.” In EJ Kontoghiorghes (ed.), *Handbook of Parallel Computing and Statistics*, pp. 481–512. Marcel Dekker/CRC Press.

Affiliation:

Péter Sólymos
 Alberta Biodiversity Monitoring Institute
 and Boreal Avian Modelling Project
 Department of Biological Sciences
 CW 405, Biological Sciences Bldg
 University of Alberta
 Edmonton, Alberta, T6G 2E9, Canada
 E-mail: solyms@ualberta.ca